

Natural Language Processing with NLTK



Vanessa Gómez Verdejo

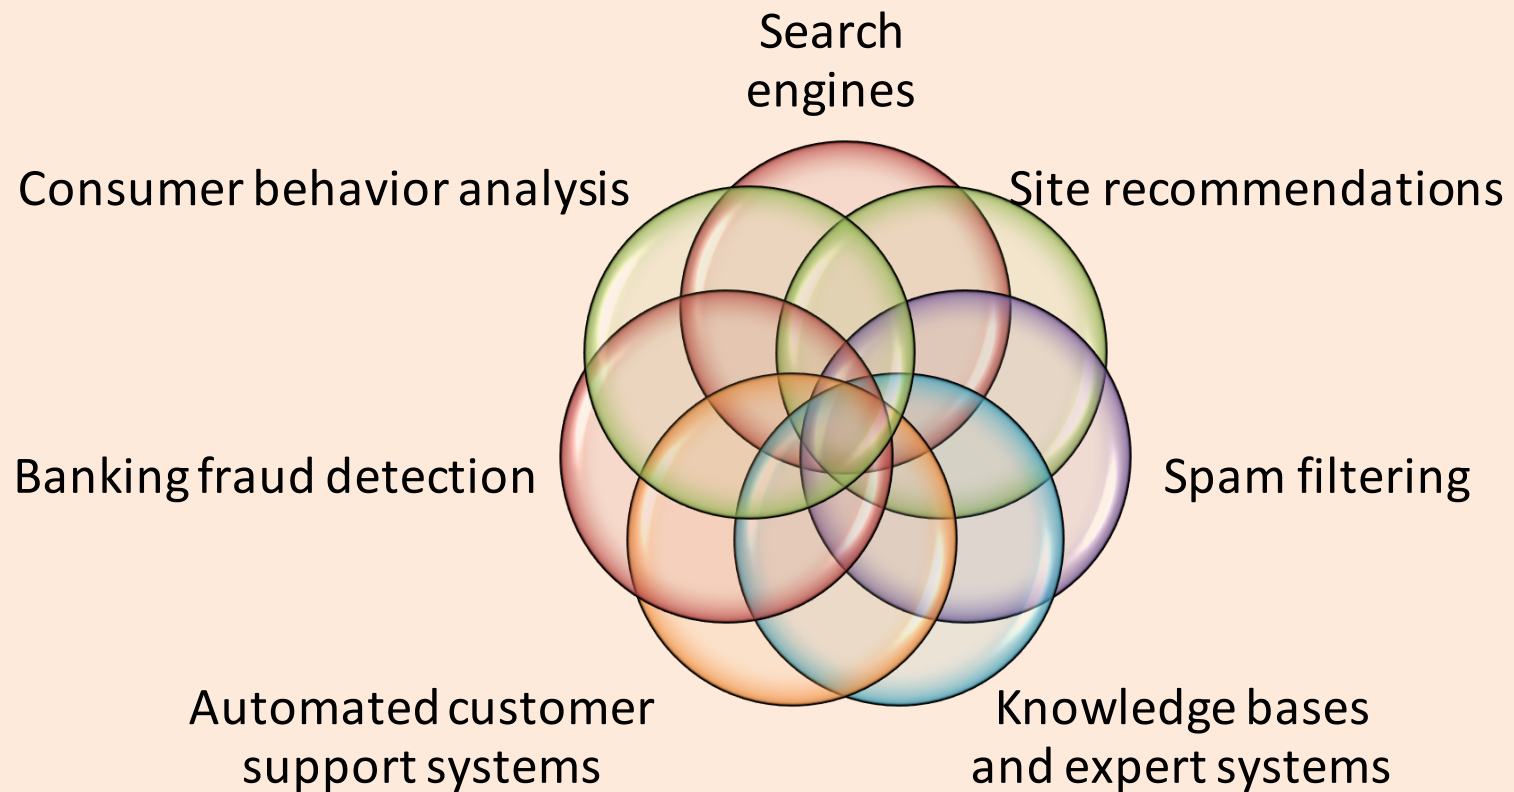
Jesús Cid Sueiro

Natural Language Processing

- Natural Language Processing:
 - Computer aided text analysis of human language.
 - The goal is to enable machines to understand human language and extract meaning from text.
 - It is a field of study which falls under the category of machine learning and more specifically computational linguistics.



Applications



Challenges to machines.

Sentiment

Ambiguity

Intent

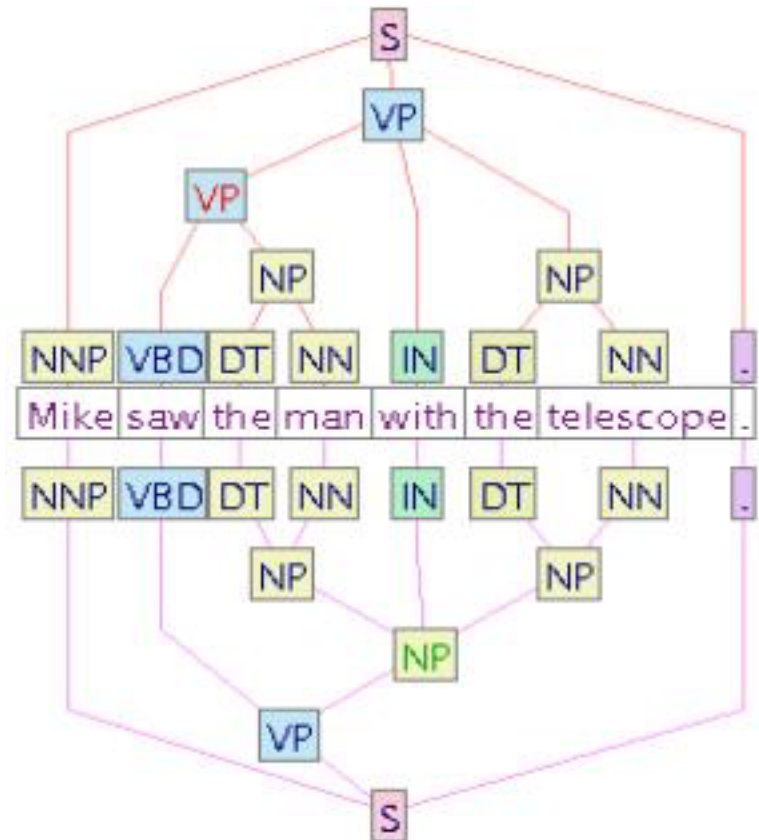
- Sarcasm
- Slang

Context

- Emphasis
- Time and date
 - Since when did “google” become a verb?

NLP with Python

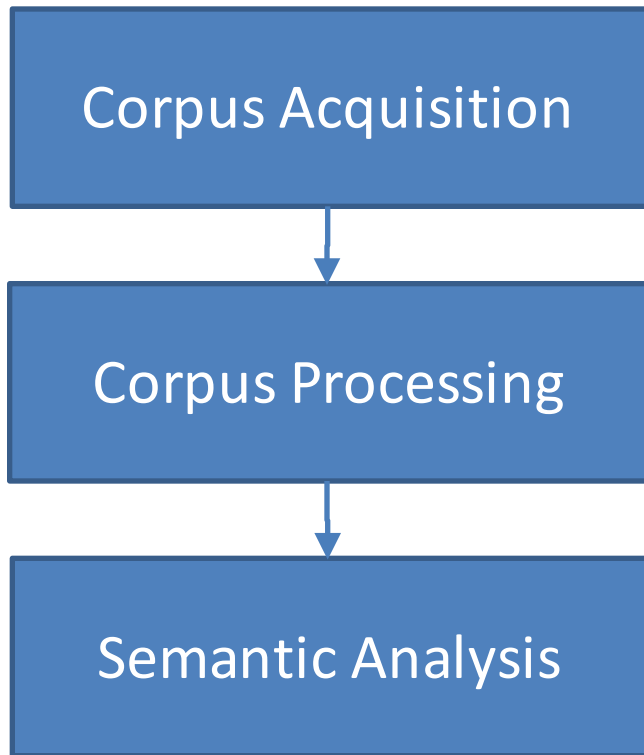
- NLTK: A package that provides
 - basic classes for representing data relevant to Natural Language Processing
 - Standard interfaces for performing NLP tasks such as tokenization, tagging and parsing
 - Standard implementation of each task which can be combined to solve complex problems
 - <http://www.nltk.org>



NLTK modules

- **corpora**: a package containing modules of example text
- **tokenize**: functions to separate text strings
- **probability**: for modeling frequency distributions and probabilistic systems
- **stem**: package of functions to stem words of text
- **wordnet**: interface to the WordNet lexical resource
- **chunk**: identify short non-nested phrases in text
- **etree**: for hierarchical structure over text
- **tag**: tagging each word with part-of-speech, sense, etc.
- **parse**: building trees over text - recursive descent, shift-reduce, probabilistic, etc.
- **cluster**: clustering algorithms
- **draw**: visualize NLP structures and processes
- **contrib**: various pieces of software from outside contributors

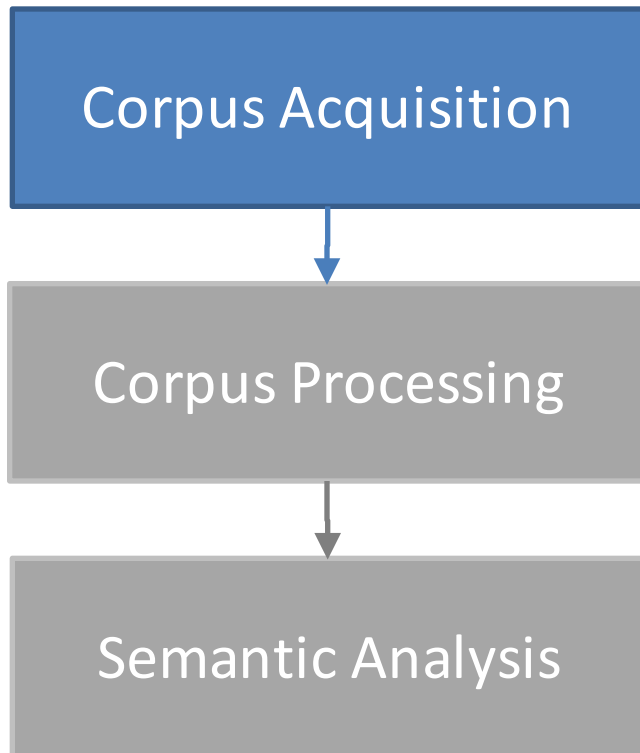
Document Corpus Analysis



Text processing tools:
Natural Language Toolkit (NLTK)

Topic Models: PLSI, LDA

Corpus Acquisition



Text processing tools:
Natural Language Toolkit (NLTK)

Topic Models: PLSI, LDA

Text sources

- Any document can be analyzed:
 - Web content: web pages, twitters, blogs, ...
 - Crawler
 - Available APIs: wikipedia
 - Local documents
 - Available corpus:
 - NLTK (see http://www.nltk.org/nltk_data/)
 - scikit-learn, gensim,...




Loading a corpus

- From NLTK (pip install nltk)

```
import nltk
nltk.download()
Mycorpus = nltk.corpus.gutenberg
text_name = Mycorpus.fileids()[0]
raw = Mycorpus.raw(text_name)
Words = Mycorpus.words(text_name)
```

Install it now and
download book content
(it takes a while)



[Emma by Jane Austen 1816]

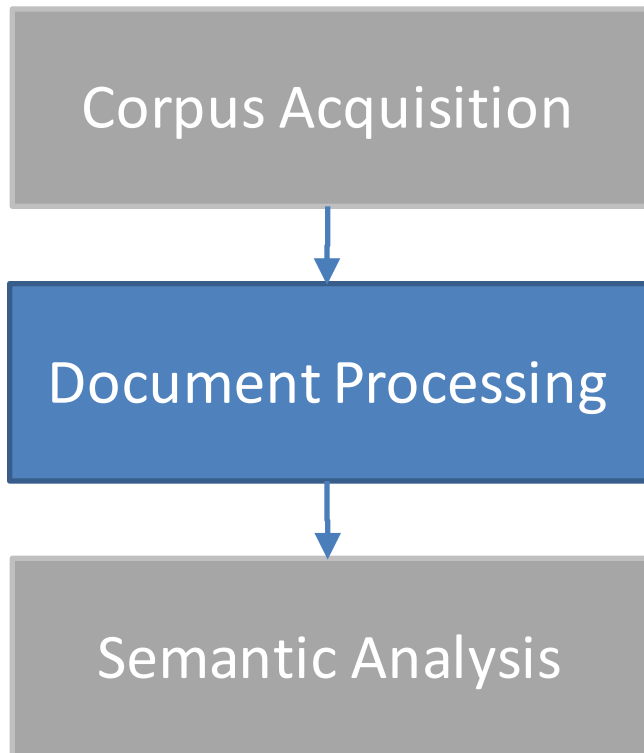
VOLUME I

CHAPTER I

Emma Woodhouse, handsome, clever, and rich, with a comfortable home and happy disposition, seemed to unite some of the best blessings of existence; and had lived nearly twenty-one years in the world with very little to distress

```
[u'[, u'Emma', u'by', u'Jane',  
u'Austen', u'1816', u']', u'VOLUME',  
u'I', u'CHAPTER', u'I', u'Emma',  
u'Woodhouse', u',', u'handsome', u',',  
u'clever', u',', u'and', u'rich',  
u',', u'with', u'a', u'comfortable',  
u'home', u'and', u'happy',  
u'disposition', u',', u'seemed',  
u'to', u'unite', u'some', u'of', ...]
```

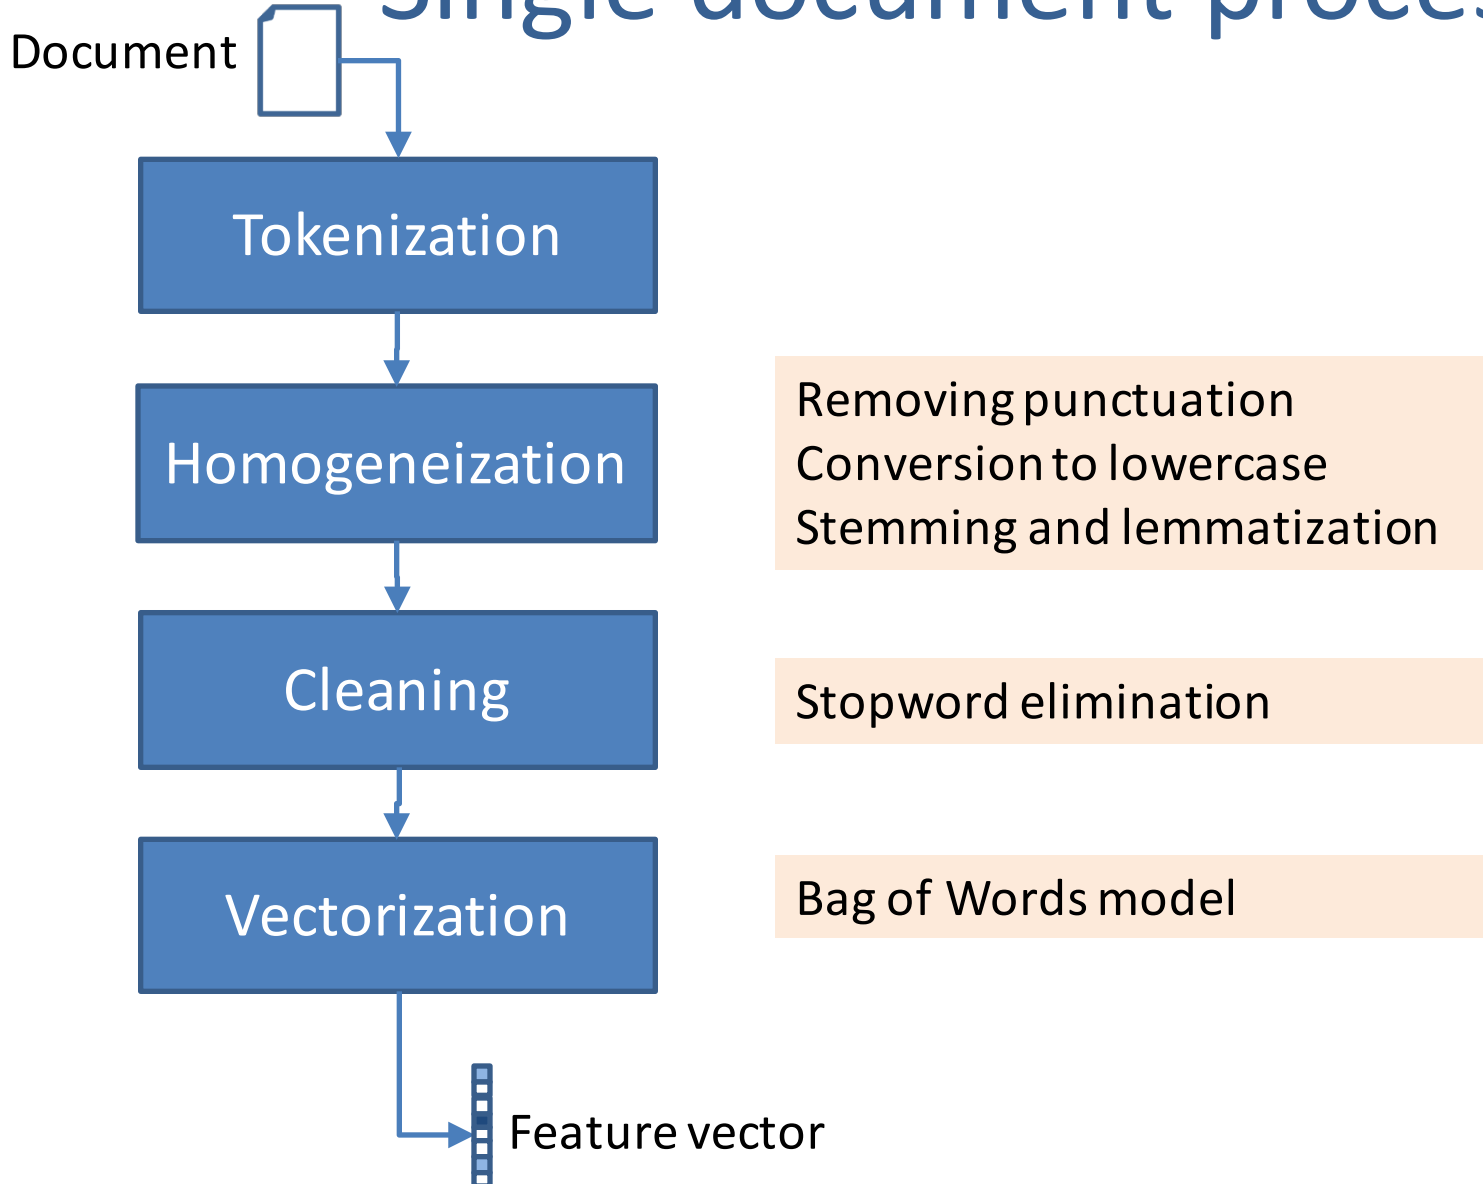
Corpus Processing



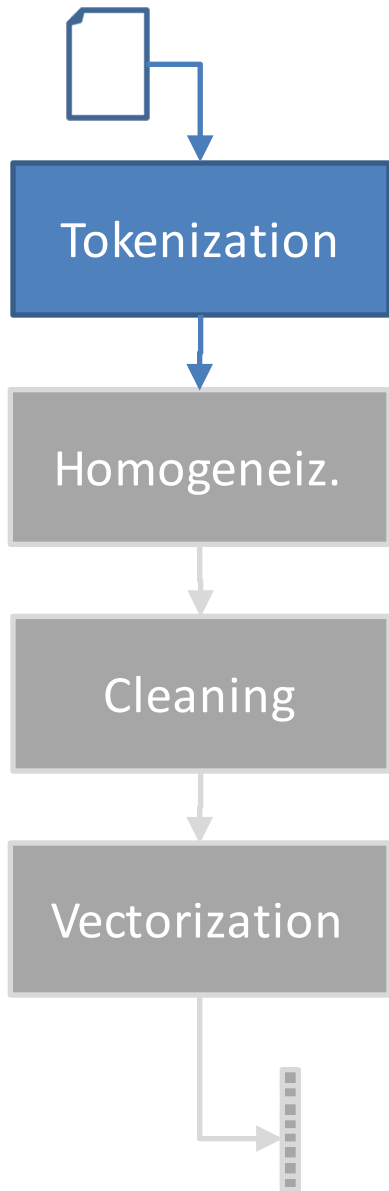
Text processing tools:
Natural Language Toolkit (NLTK)

Topic Models: PLSI, LDA

Single document processing



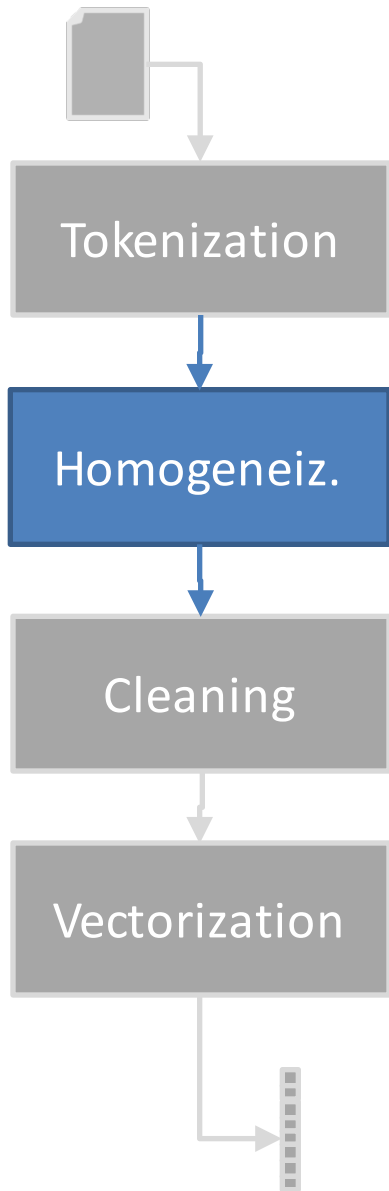
Tokenization



- From text to words (elements inside a sentence):

```
>> Words = corpus.words(text_name)
>> sentence = "Hola, mundo."
>> sentence.split()
['Hola,', 'mundo.']
>> from nltk.tokenize import word_tokenize
>> word_tokenize(sentence)
['Hola', ',', 'mundo', '.']
```

Homogeneization



- EXERCISE 1

- Convert every word to lowercase

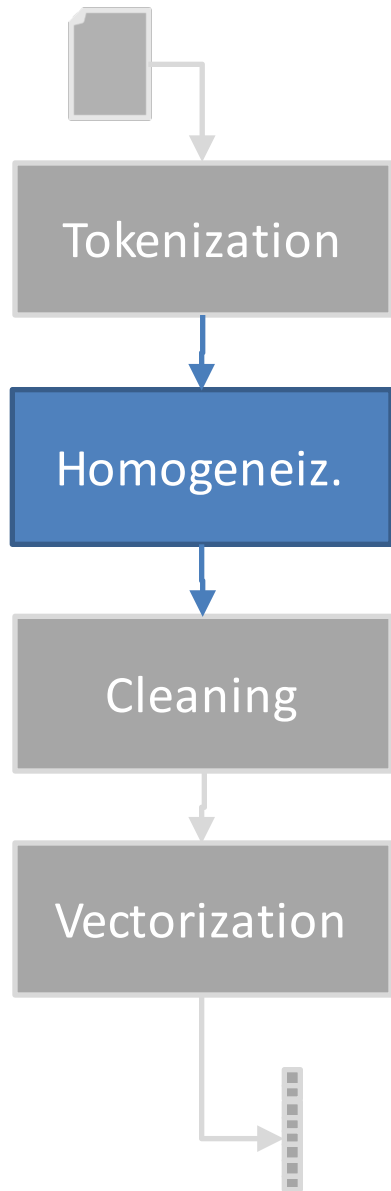
```
>> clean_text = [w.lower() for w in text]
```

- EXERCISE 2

- Remove punctuation

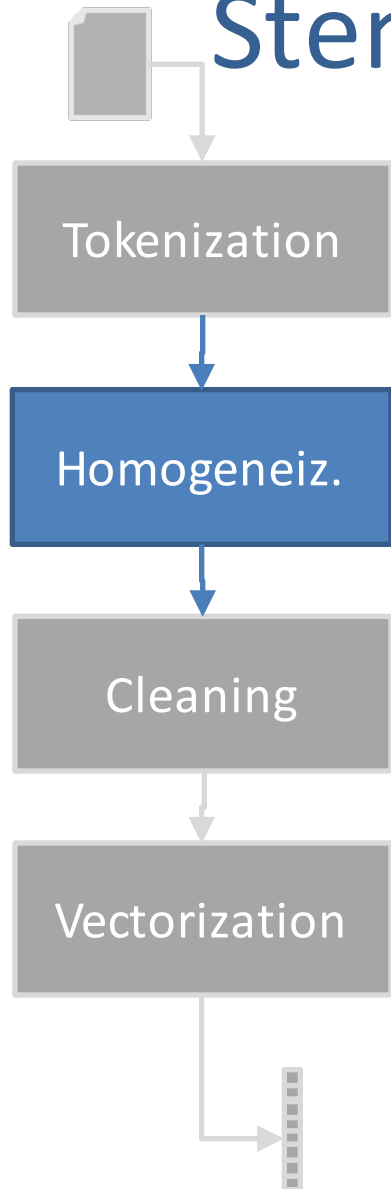
```
>> clean_text = [w for w in text  
                  if w.isalnum()]
```

Homogeneization



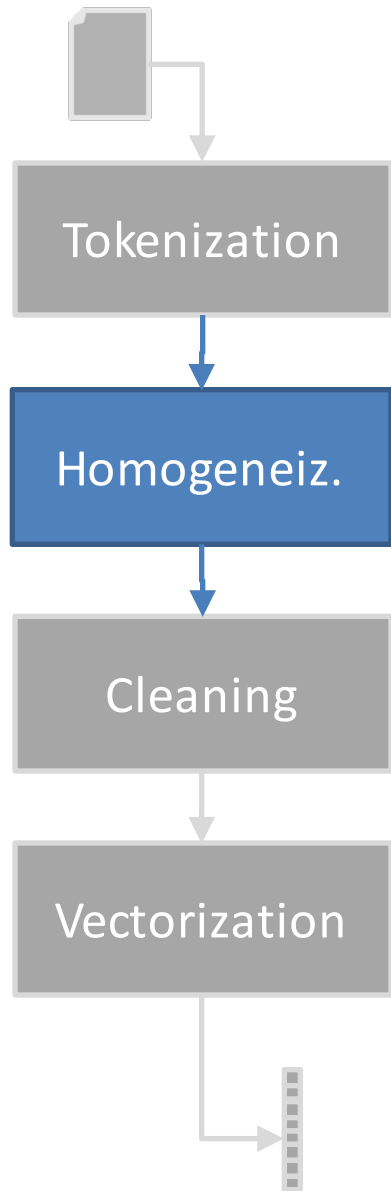
- The class string in python:
 - `s.find(t)` index of first instance of string `t` inside `s` (-1 if not found)
 - `s.rfind(t)` index of last instance of string `t` inside `s` (-1 if not found)
 - `s.join(text)` combine the words of the text into a string using `s` as the glue
 - `s.split(t)` split `s` into a list wherever a `t` is found (whitespace by default)
 - `s.lower()` a lowercased version of string `s`
 - `s.upper()` an uppercased version of string `s`
 - `s.title()` a titlecased version of string `s`
 - `s.strip()` a copy of `s` without leading or trailing whitespace
 - `s.replace(t, u)` replace instances of `t` with `u` inside `s`
 - `t in s` test if `t` is contained inside `s`

Stemming and Lemmatization



- Motivation:
 - Different forms of a word
 - organize, organizes, organizing.
 - Derivationally related words with similar meanings:
 - democracy, democratic, democratization.
- Goal: to reduce inflectional or derivationally related forms of a word to a common base form.
 - am, are, is → be
 - car, cars, car's, cars' → car
- Stemming:
 - Chops off the ends of words in the hope of achieving this goal correctly most of the time.
 - See, saw → s
- Lemmatization:
 - Usually refers to doing things properly with a vocabulary and morphological analysis of words, aiming to return the base or dictionary form (lemma) of a word.
 - 'saw' → 'see' if verb, 'saw' if noun.
- Stemming vs lemmatization
 - Stemming commonly collapses derivationally related words, whereas lemmatization commonly only collapses the different inflectional forms of a lemma.


Stemming



- We count similar words in different variants as different words
- We need a function that reduces words to their specific word stem.

```
import nltk.stem  
s = nltk.stem.SnowballStemmer('english')  
s.stem("imaging")    # → u'imag'  
s.stem("image")      # → u'imag'
```

Lemmatization

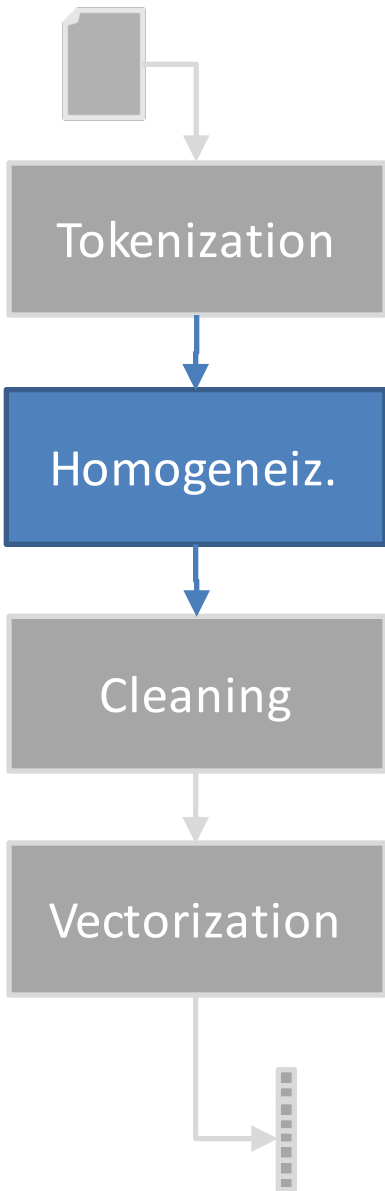


- Lemmatization is a more complex process.
- Lemmatization in NLTK uses WordNet.

```
from nltk.stem import WordNetLemmatizer
>> wnl = WordNetLemmatizer()
>> print(wnl.lemmatize('dogs'))
dog
>> print(wnl.lemmatize('churches'))
church
>> print(wnl.lemmatize('abaci'))
abacus
```

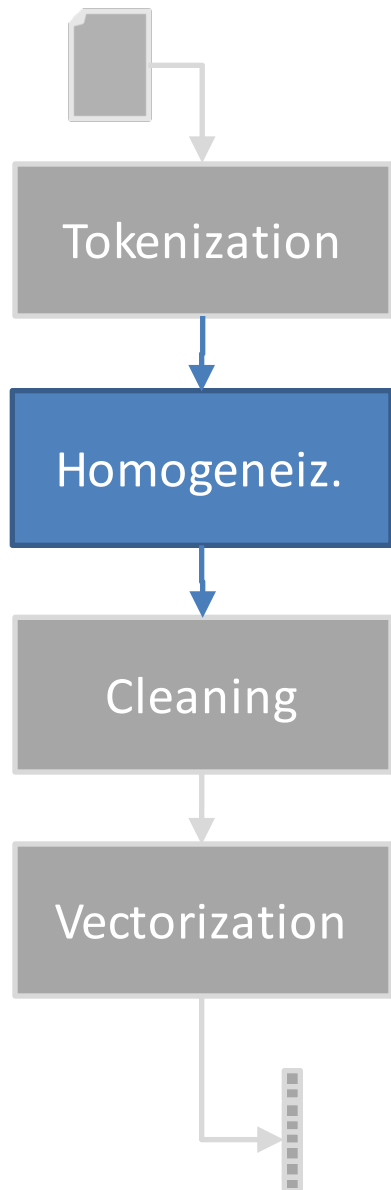
Lemmatization

- With contextual information (the grammatical role of the word) `.lemmatize()` can filter grammatical differences.



```
from nltk.stem import WordNetLemmatizer
>> wnl = WordNetLemmatizer()
>> print(wnl.lemmatize('is'))
is
>> print(wnl.lemmatize('is', pos='v'))
be
```

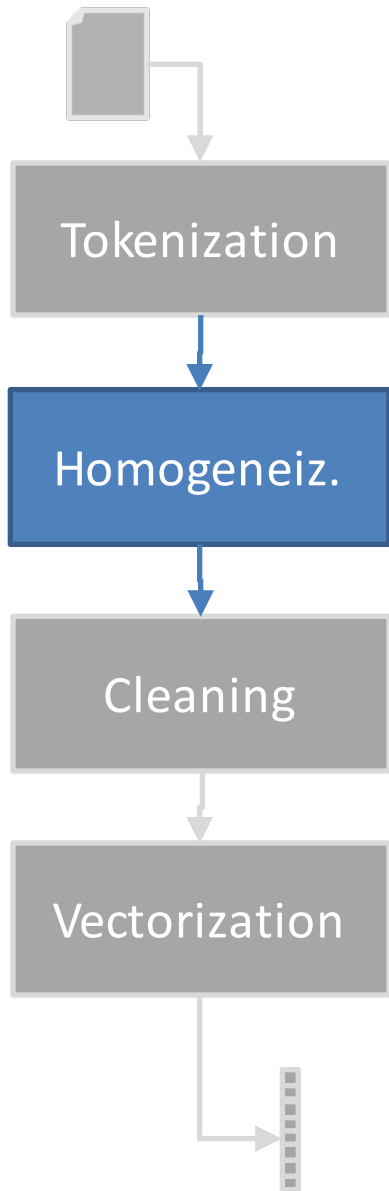
Lemmatization



- Part of Speech Tagging.
 - Part-of-speech (POS) tagging is the process of assigning a word to its grammatical category (noun, verb, adverb,...), in order to understand its role within the sentence.
 - POS taggers typically take a sequence of words (i.e. a sentence) as input, and provide a list of tuples (word, pos) as output.
 - POS tagging is what provides the contextual information to `.lemmatize()` to filter grammatical differences.

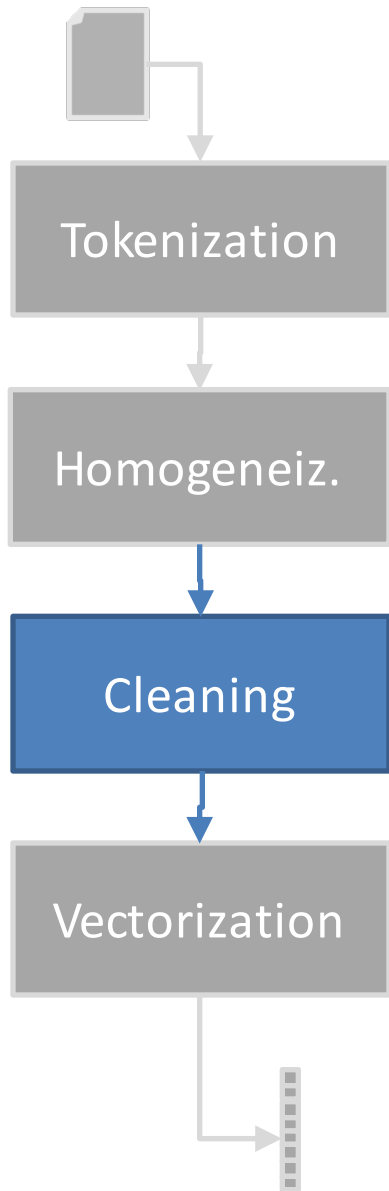
```
>> from nltk import pos_tag
>> s = "This is a simple sentence"
>> tokens = word_tokenize(s)
>> tokens_pos = pos_tag(tokens)
>> print(tokens_pos)
[('This', 'DT'), ('is', 'VBZ'), ('a', 'DT'),
 ('simple', 'JJ'), ('sentence', 'NN')]
```

Homogeneization




- *N*-grams
 - Some words tend to occur in groups
 - information processing, machine learning...
 - It can be useful that they are analyzed in groups
 - There are routines to detect them, but the easiest way is...
 - information processing
 - machine learning

Cleaning




- Removing least relevant words
 - Some words appear very often in all sorts of different contexts.
 - They are so frequent that they do not help to distinguish between different texts.
 - These words are called stop words.
 - The best option would be to remove them

Cleaning

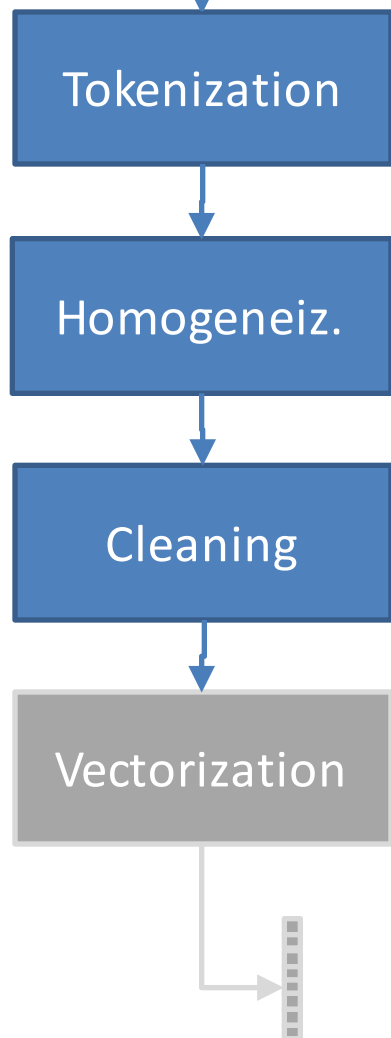


- Removing stopwords:

```
from nltk.corpus import stopwords  
sw = stopwords.words('english')  
clean_text = [word for word in document  
               if not word in sw]
```




Parallel Document Processing

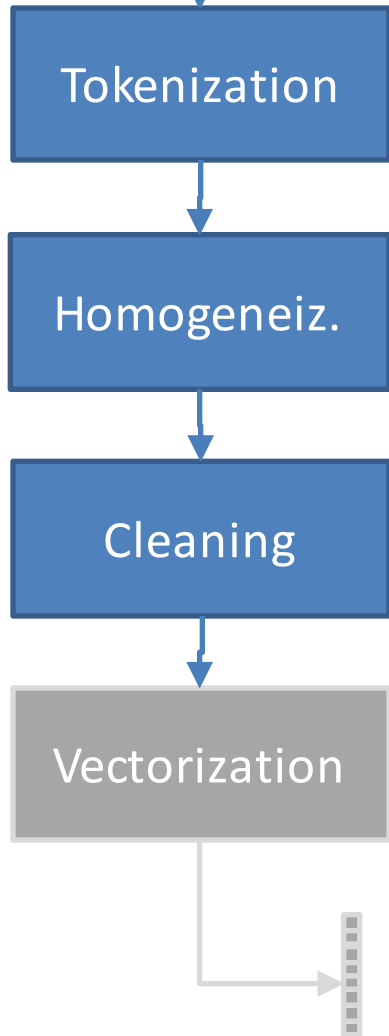


- Until now, we have worked with a single document
- Extend your code to work with all the documents of the corpus
- Create a list of text, where each row is a previously processed text

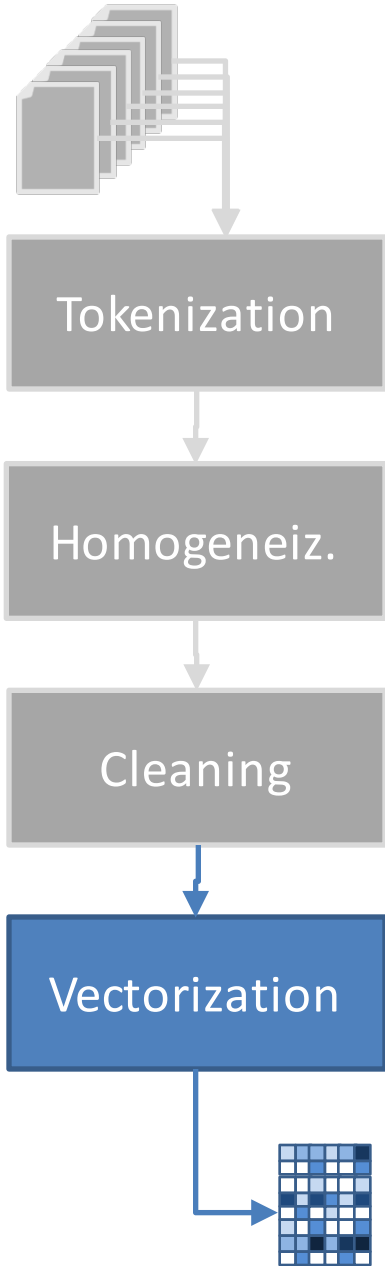
```
content = [  
    [u'fulton', u'counti', u'grand', ..., u'said', u'friday']  
    [u'austin', u'texa', u'committe', ..., u'price', u'abandon']  
    ....  
    [u'dear', u'sir', u'let', u'begin', ..., u'mind', u'address']  
]
```

Parallel Document Processing



```
content = []  
for text_name in corpus.fileids():  
    path = nltk.data.find(  
        'corpora/brown/'+text_name)  
    f = open(path, 'rU')  
    raw = f.read()  
    # Here you can process your  
    # raw text → clean_text  
    content.append(clean_text)  
    f.close()
```





Vectorization

- Bag of words: counting words
 - ML algorithms process numbers, not words.
 - Only if we manage to transform text into meaningful numbers, we can feed it into ML algorithms
 - Bag-of-word approach: for each word in the document, count its occurrence and note it in a vector

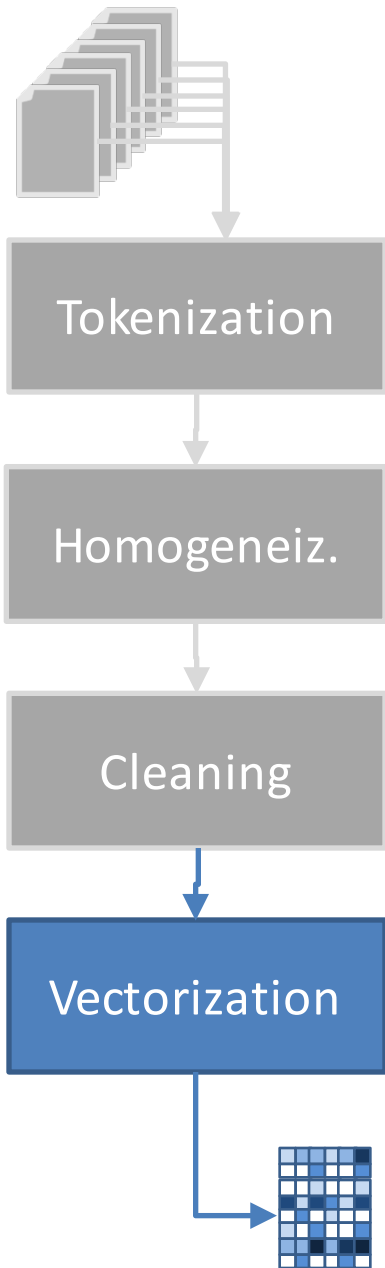
Vectorization

- Bag of words: counting words



Vectorization

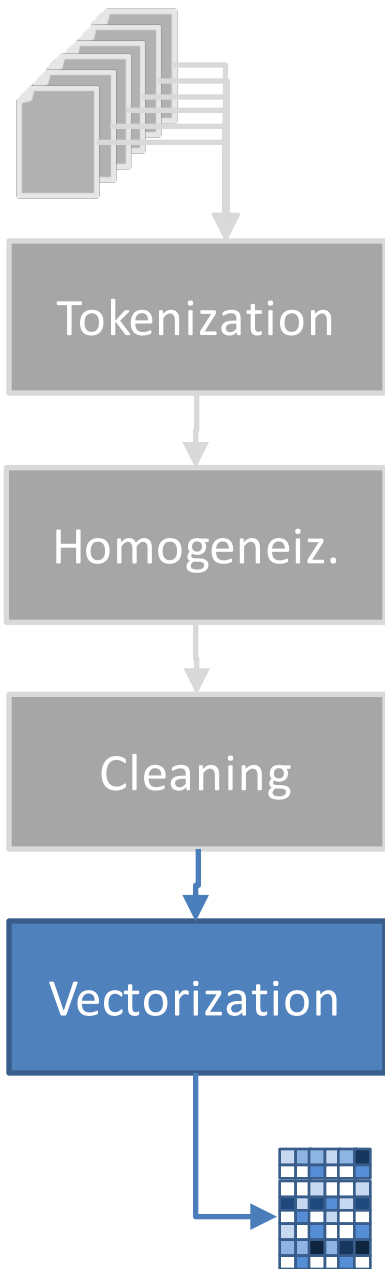
- Bag of words: counting words



[Hadoop, is, great, python, is, great]

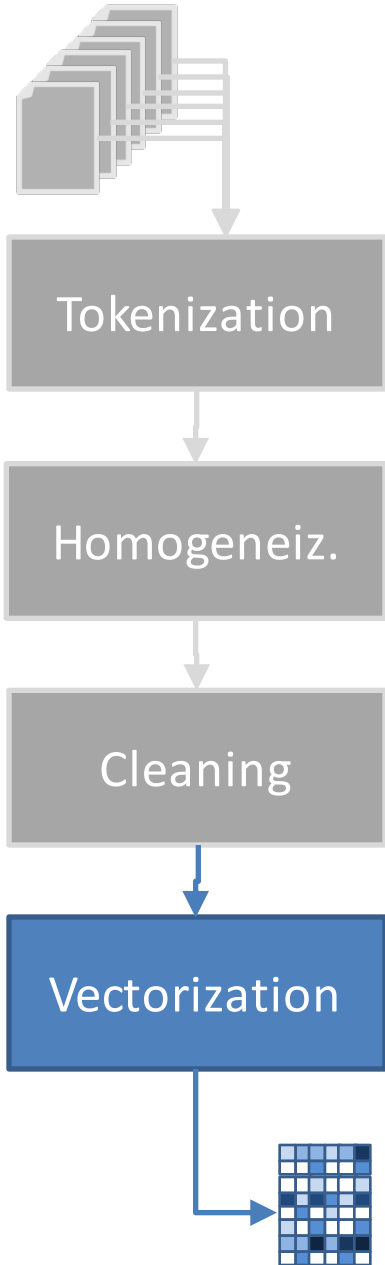
[(Hadoop, 1), (is, 2), (great, 2), (python, 1)]

[(12342, 1), (3423, 2), (5676, 2), (6768, 1)]



Vectorization

- Term frequency - Inverse document frequency (TF-IDF)
 - BoW: the feature values simply count occurrences of terms in a document.
 - High occurrence terms?? They appear in all documents → NON RELEVANT.
 - Low occurrence terms?? They appear in very few documents → RELEVANT
 - This can only be solved by:
 - counting term frequencies for each document
 - discounting those that appear in many posts



Vectorization

- Term frequency - Inverse document frequency (TF-IDF)
 - We want a high value for a given term in a given doc if that term occurs often in that particular doc and very rarely anywhere else
 - $TF(w, d) = \frac{bow(w, d)}{\# \text{ words in doc}}$
 - $IDF(w, d) = \log \frac{\# \text{ docs}}{\# \text{ docs with } w}$
 - $TF-IDF(w, d) = TF(w, d) \times IDF(w, d)$
 - IDF \rightarrow 0 in common words & IDF increases in rare words